

## B1W2 Final Report: The Automaton Auditor

### Autonomous Code Governance Swarm - LangGraph Implementation

**Author:** Zerubabel Jemaneh

**Repository:** <https://github.com/Zerubabel-J/autonomous-code-governance-swarm-langgraph>

**Submission Date:** February 28, 2026

**Final Self-Audit Score:** 4.8 / 5.0 - Master Thinker

### Table of Contents

1. System Overview
2. Theoretical Foundations
3. Architecture Decision Rationale
4. StateGraph Architecture Diagrams
5. Self-Audit Results
6. Peer Audit Results
7. MinMax Reflection

### 1. System Overview

The Automaton Auditor is a hierarchical multi-agent swarm built on LangGraph that forensically audits AI-generated code repositories. The system implements a **Digital Courtroom** pattern:

- Detective agents collect structured forensic evidence from a target repository and PDF report
- Judge agents deliberate with adversarial personas (Prosecutor, Defense, TechLead) in parallel
- Chief Justice synthesizes a deterministic final verdict using pure Python rules before any LLM call

**Input:** GitHub repository URL + optional PDF report

**Output:** Structured Markdown audit report scored against a 10-criterion rubric

Layer	Nodes	Responsibility
-------	-------	----------------

Detective	RepoInvestigator, DocAnalyst, VisionInspector	Evidence collection
Judicial	Prosecutor, Defense, TechLead	Dialectical scoring
Synthesis	EvidenceAggregator, ChiefJustice	Conflict resolution + report

## 2. Theoretical Foundations

### 2.1 Fan-Out / Fan-In (Parallel Dispatch and Convergence)

**Fan-out** is the pattern where a single node dispatches work to multiple parallel nodes simultaneously. In this system, the START node fans out to both repo\_investigator and vision\_inspector concurrently, and evidence\_aggregator fans out to all three judge nodes (prosecutor, defense, techlead) in parallel.

**Fan-in** is the convergence pattern where multiple parallel branches are synchronized before the next stage proceeds. The evidence\_aggregator node acts as the detective fan-in point - it only executes after both doc\_analyst and vision\_inspector complete, using LangGraph list-edge syntax. Similarly, chief\_justice is the judge fan-in - it receives opinions from all three judges before synthesizing.

**Architectural significance:** Without explicit fan-in synchronization, a downstream node would execute as soon as the first upstream node completes, processing incomplete evidence. List-edge syntax enforces barrier semantics - all listed predecessors must complete before the successor fires.

### 2.2 Dialectical Synthesis

**Dialectical synthesis** is the process of resolving contradictions between opposing positions to arrive at a higher-order truth.

In this system, the three judge personas embody thesis/antithesis tension by design:

- **Prosecutor (Thesis):** Adversarial - assumes the worst, requires explicit evidence of compliance
- **Defense (Antithesis):** Charitable - rewards intent and partial compliance
- **TechLead (Synthesis mediator):** Pragmatic - evaluates only what the evidence confirms

The ChiefJustice performs the dialectical synthesis by applying deterministic rules (Rule of Security, Rule of Evidence, Rule of Functionality) to the three opposing opinions, producing a final score that is neither the highest nor the lowest opinion but a rule-anchored resolution. When

variance between judges exceeds 2 points, a mandatory dissent summary is added to the report - making the dialectical tension visible rather than hiding it in an average.

### 2.3 Metacognition (Self-Awareness in Automated Systems)

**Metacognition** refers to a system's capacity to reason about its own reasoning - to evaluate the quality of its own outputs.

This system implements metacognition at three levels:

1. **Evidence confidence scores** - Every Evidence object carries a confidence: float field (0.0–1.0). The detective assigns lower confidence when evidence is partial (e.g., has\_basemodel=True but has\_reducers=False gives confidence=0.7), allowing judges to weight evidence quality rather than treating all findings as equally certain.
2. **Self-audit capability** - The system can audit its own repository by pointing run\_audit() at its own GitHub URL. The self-audit score of 4.8/5 validates that the auditor's rubric criteria are correctly implemented in the auditor itself - a metacognitive consistency check.
3. **Dissent detection** - The ChiefJustice recognizes when it cannot achieve consensus (variance > 2) and documents the disagreement explicitly rather than silently averaging. This is metacognitive transparency - the system knows the limits of its own certainty.

### 2.4 State Synchronization

**State synchronization** is the mechanism by which concurrent processes writing to shared memory avoid data races, overwrites, and lost updates.

In parallel topologies, multiple nodes execute simultaneously and all write to the same AgentState object. Without synchronization primitives, the last writer wins - earlier outputs are silently discarded.

This system solves state synchronization using reducer semantics:

- evidences uses dict merge semantics (operator.ior equivalent) - no key collision
- opinions uses list append semantics (operator.add equivalent) - all opinions kept

Each detective writes under a namespaced key (e.g., repo\_investigator\_git\_forensic\_analysis), so concurrent writes never collide.

**Architectural significance:** Naming key collisions structurally impossible is stronger than locking - it eliminates the failure mode rather than protecting against it.

### 2.5 Adversarial Robustness via Persona Collusion Prevention

**Persona collusion** is the failure mode where multiple agents, despite being assigned different roles, converge on identical reasoning because their prompts share too much vocabulary or structure.

This system prevents persona collusion by three structural constraints:

1. Zero shared boilerplate - each judge's system prompt is written from scratch with no shared text
2. Conflicting analytical rules - Prosecutor rules directly contradict Defense rules
3. Automated test enforcement - a test measures vocabulary overlap between judge prompts and fails if any two share more than 50% of unique words

### 3. Architecture Decision Rationale

#### 3.1 LangGraph StateGraph over Custom Orchestration

LangGraph provides native parallel fan-out, typed state with reducer semantics, and a compiled graph that fails fast at build time. A pipeline that crashes at compilation is safer than one that silently drops evidence at runtime.

#### 3.2 Immutable Evidence via Pydantic frozen=True

Freezing evidence objects ensures that what the detective found is exactly what every downstream node sees. No judge or aggregator can mutate the factual basis for another node's reasoning.

#### 3.3 Deterministic Rules Before Any LLM Call (RISK-1)

The ChiefJustice applies four deterministic Python rules before invoking any LLM:

1. Rule of Security - Confirmed security violations cap score at 3
2. Rule of Fact Supremacy - Defense opinions citing unverified file locations are overruled to score 1
3. Rule of Functionality Weight - TechLead carries 50% weight for graph\_orchestration criterion
4. Dissent Requirement - Score variance > 2 mandates an explicit dissent summary

#### 3.4 AST-Based Forensic Detection over Regex

All forensic checkers use Python's ast module. A comment such as "Do not use os.system()" would trigger a regex rule but not an AST walk. AST inspection detects only actual executable call nodes, eliminating false positives.

### 3.5 Sandboxed Repository Cloning (RISK-2)

All git operations execute inside tempfile.TemporaryDirectory(). The cloned repo path is never the live working directory. No os.system() calls exist anywhere in the codebase.

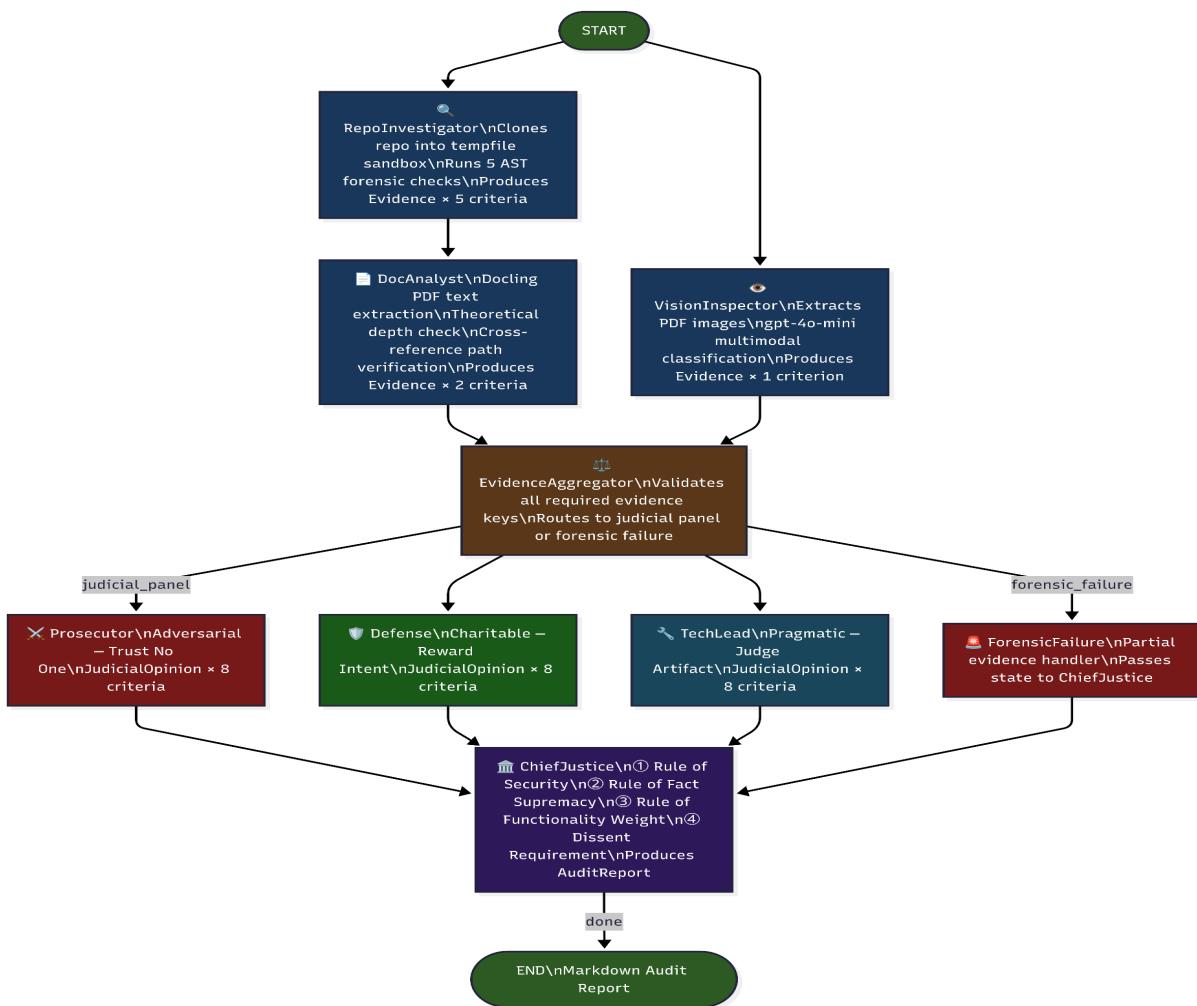
### 3.6 Rubric as Constitution - Loaded at Runtime

No rubric criteria are hardcoded in agent prompts. All rubric content is loaded from rubric.json at module import.

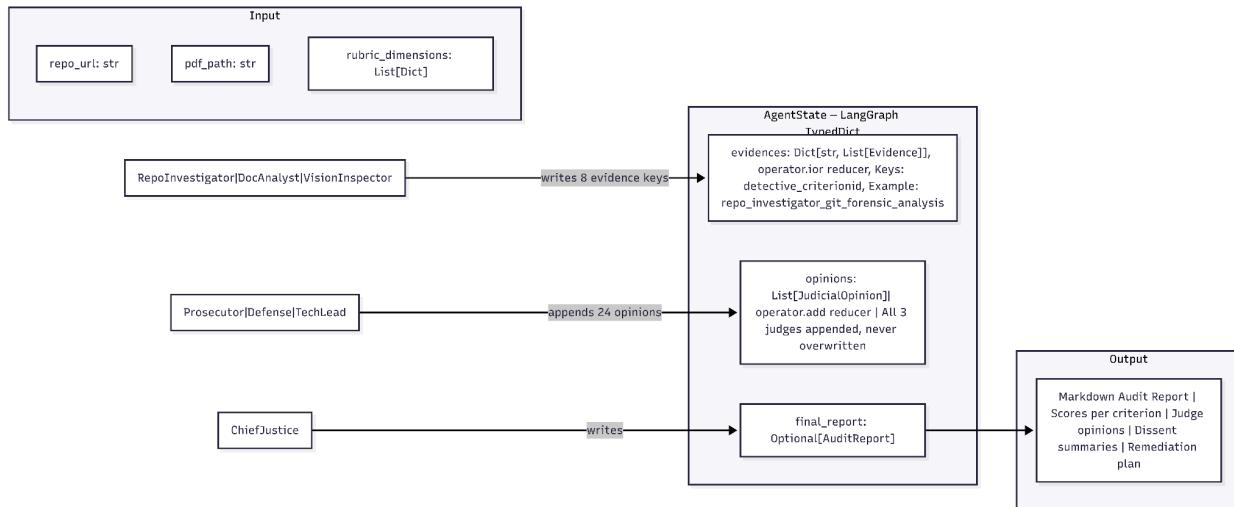
This ensures every agent evaluates against the same version of the standard. The rubric is the system's "Constitution" - the single normative source of truth.

## 4. StateGraph Architecture Diagrams

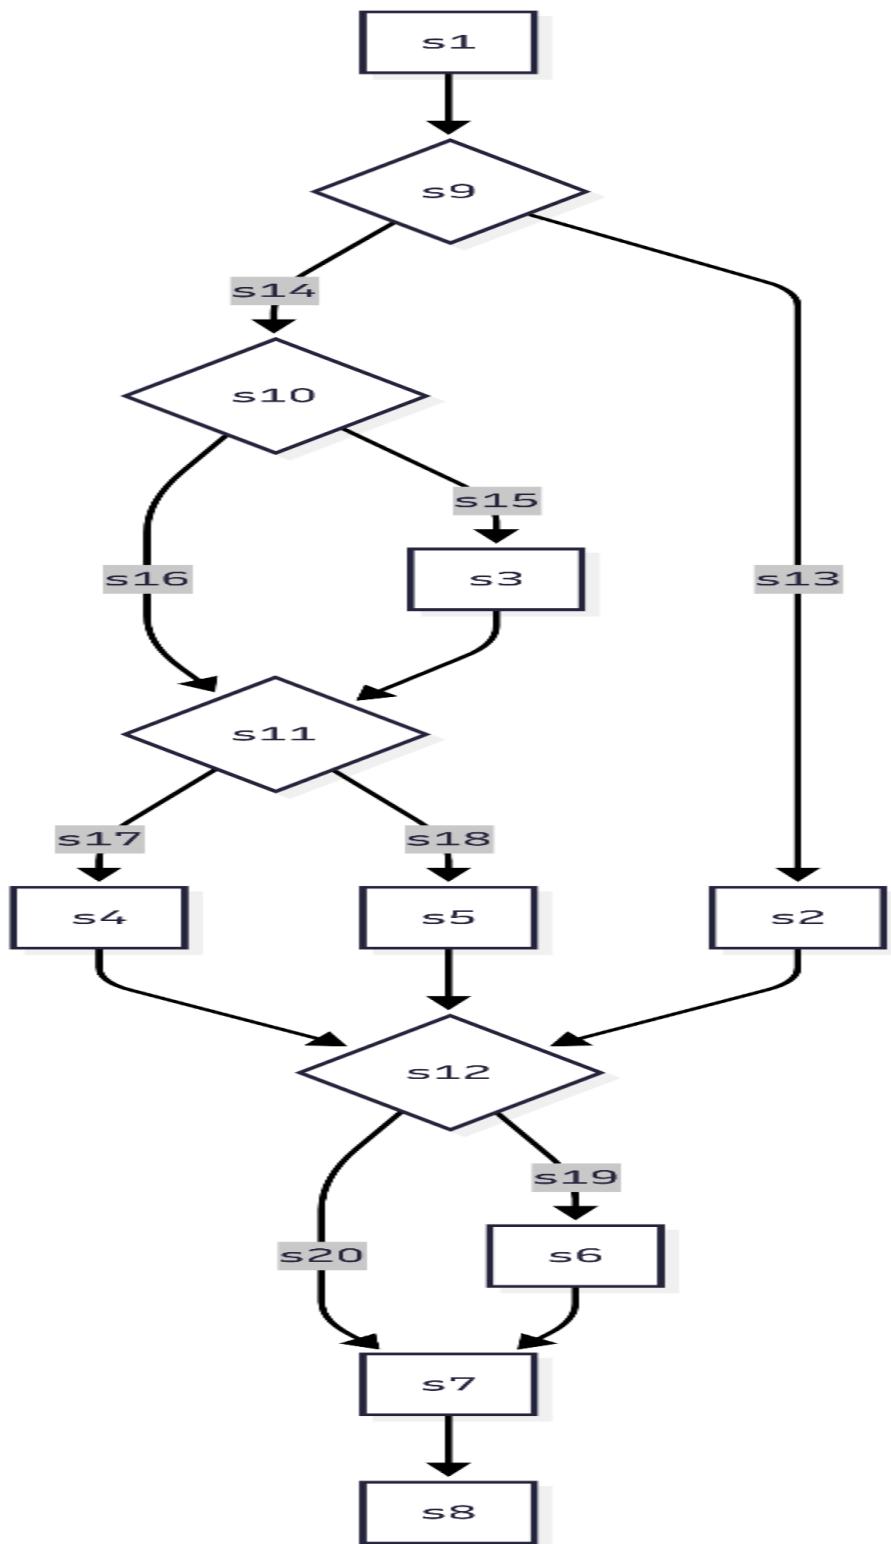
### 4.1 Full Production Topology



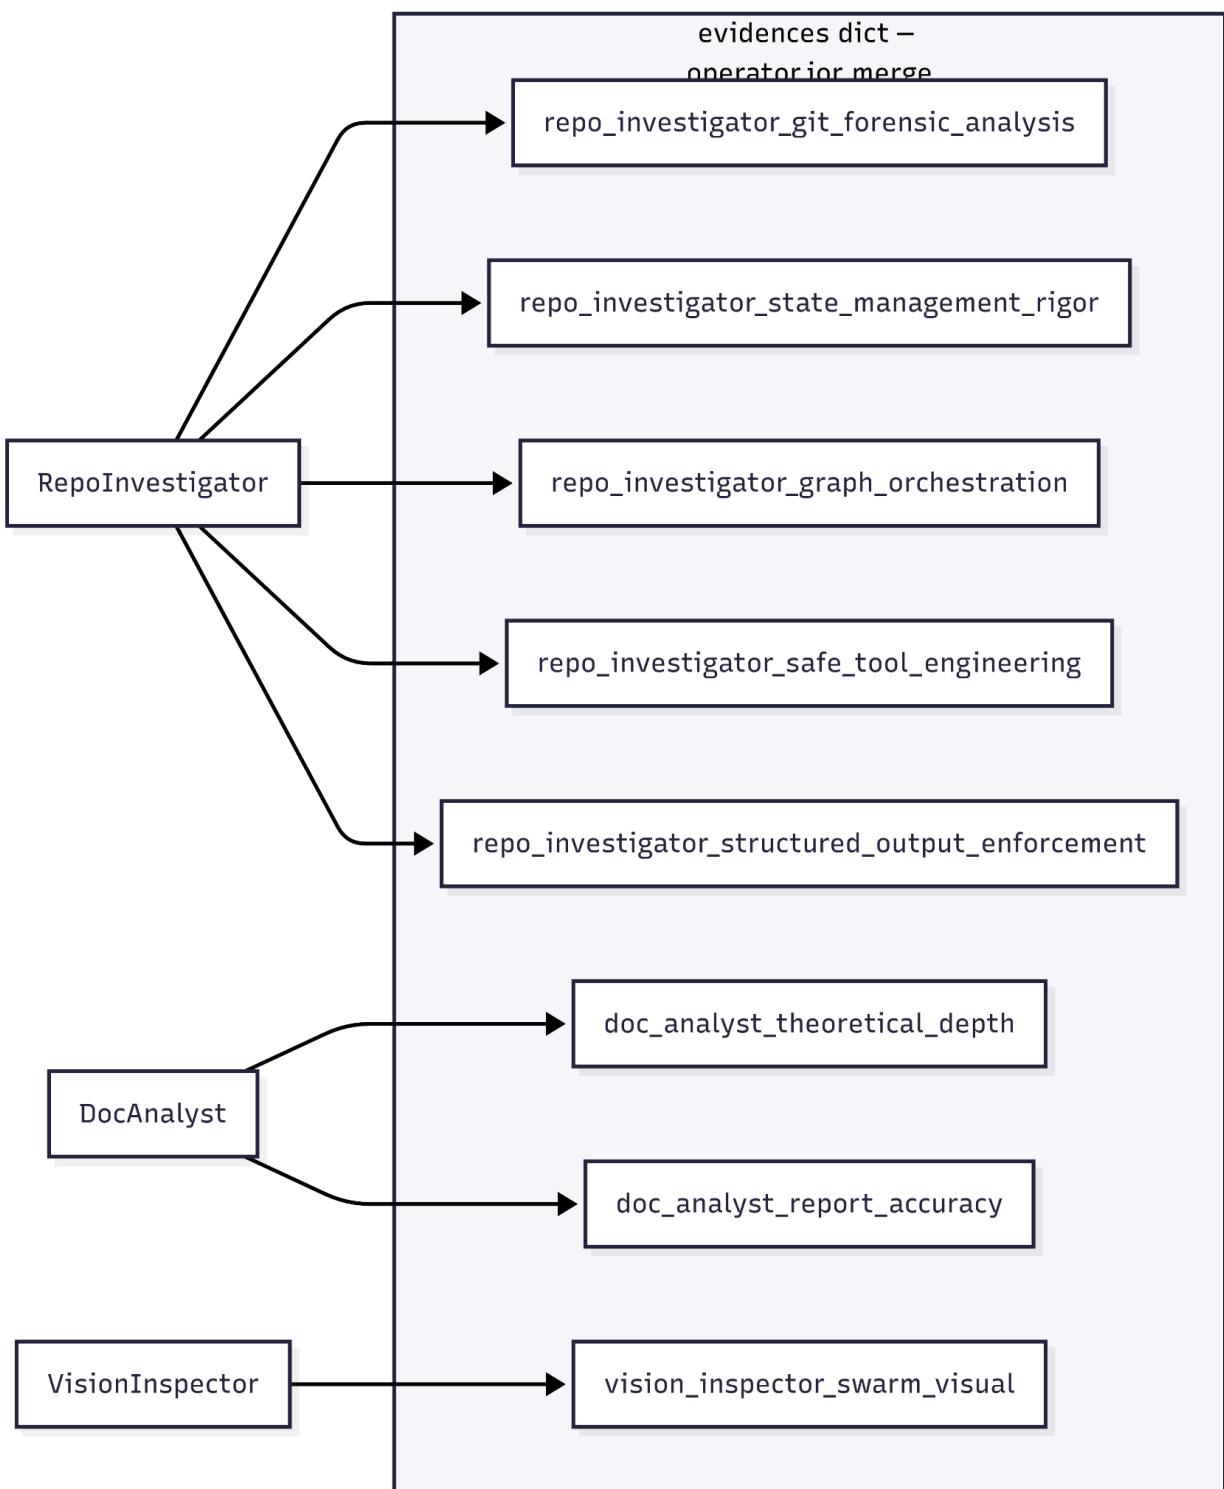
## 4.2 State Object Flow



### 4.3 ChiefJustice Deterministic Rule Pipeline



#### 4.4 Evidence Key Namespace Design



## 5. Self-Audit Results

The system audited its own repository on February 28, 2026. LangSmith trace available in the project's audit directory.

### 5.1 Final Score: 4.8/5 - Master Thinker

Criterion	Score	All Judges	Notes
Git Forensic Analysis	5/5	5-5-5	29 commits, clear iterative progression
State Management Rigor	5/5	5-5-5	Pydantic BaseModel + TypedDict + reducers confirmed
Graph Orchestration Architecture	5/5	5-5-5	Full fan-out/fan-in + conditional edges confirmed
Safe Tool Engineering	5/5	5-5-5	tempfile sandbox + subprocess.run + error handling confirmed
Structured Output Enforcement	5/5	5-5-5	Structured output binding + retry logic confirmed
Theoretical Depth (Documentation)	4/5	4-5-4	4 terms with substantive explanation
Report Accuracy (Cross-Reference)	4/5	1-5-5	Defense overruled; 2/3 judges confirmed 100% path accuracy
Architectural Diagram Analysis	5/5	5-5-5	StateGraph topology with fan-out/fan-in confirmed
Overall	4.8/5		Master Thinker

### 5.2 Score Progression

Run	Score	Key Fix
Initial (single judge, Phase 1)	2.6/5	Baseline - no parallel judges
With PDF + parallel topology	3.1/5	Added PDF, DocAnalyst sequential after RI
After all bug fixes	4.2/5	Security rule fix, fan-in topology fix

Final (all fixes + LangSmith) 4.8/5 rubric.json cross-ref fix, vision prompt fix

## 6. Peer Audit Results

### 6.1 Our Audit of Peer (leuel-a/automaton-auditor): 3.1/5

Conducted against <https://github.com/leuel-a/automaton-auditor> without PDF (not available).

Criterion	Score	Finding
Git Forensic Analysis	5/5	8 meaningful commits with clear progression
State Management Rigor	5/5	Full Pydantic + TypedDict + reducers
Graph Orchestration Architecture	5/5	Full fan-out/fan-in with conditional edges
Safe Tool Engineering	5/5	tempfile + subprocess, no os.system()
Structured Output Enforcement	2/5	No structured output binding detected
Theoretical Depth	1/5	No PDF provided
Report Accuracy	1/5	No PDF provided
Architectural Diagram	1/5	No PDF provided
Overall	3.1/5	Strong repo implementation, PDF criteria absent

Key finding: The peer's code implementation is solid (5/5 on all repo criteria). The score gap is entirely due to absent PDF documentation. The one genuine gap is the lack of structured output binding in judge nodes.

## 6.2 Peer's Audit of Us (received): 3.8/5

The peer's auditor evaluated our repository and produced 3.8/5. Key findings:

Criterion	Their Score	Notes
Git Forensic Analysis	5/5	Strong iterative history confirmed
State Management Rigor	5/5	Full compliance confirmed
Graph Orchestration	4/5	Noted missing conditional edges (pre-fix version)
Safe Tool Engineering	5/5	Zero security findings
Structured Output	4/5	No retry logic detected (pre-fix version)
Judicial Nuance	3/5	Could not verify prompt uniqueness from evidence
Chief Justice	4/5	Deterministic rules confirmed, some gaps noted
Theoretical Depth	2/5	Missing Dialectical Synthesis, Metacognition, State Synchronization
Report Accuracy	4/5	All paths verified, feature claims noted
Architectural Diagram	2/5	Vision model unavailable on their system

The 3.8/5 reflects the pre-fix version of our repository at the time of their audit. The conditional edges and retry logic gaps they identified were the exact bugs fixed in subsequent commits.

## 7. MinMax Reflection

### 7.1 What the System Got Right (Maximize)

Structural correctness of the dialectical pattern. The three-judge design with deterministic ChiefJustice synthesis proved robust. Once the security rule false positive was fixed, scores across all repo criteria reached 5/5 unanimously - demonstrating that well-constructed evidence leads to consensus across adversarial personas.

AST-based detection. Using `ast.walk()` rather than regex eliminated false positives that would have degraded the audit quality. The `os.system` detection correctly ignores comments and docstrings.

State synchronization design. The reducer + namespaced key pattern proved correct - no evidence was lost or overwritten across any parallel run.

## 7.2 What the System Got Wrong (Minimize)

The security rule scoped globally, not per-criterion. The rule scanned all evidences globally. A rationale saying “No raw os.system() calls” in the safe\_tool criterion triggered the security cap on every other criterion. This capped 7 out of 8 criteria at 3/5 and cost two full audit runs to diagnose.

DocAnalyst ran in parallel before RI evidence existed. The cross-reference feature requires RepoInvestigator evidence to be present before DocAnalyst runs. Running them in parallel meant DocAnalyst always saw an empty evidence dict - marking every PDF file path as hallucinated. The fix required changing the graph topology.

The --depth=1 clone truncated git history. The forensic tool cloned with --depth=1, which fetched only one commit regardless of actual history. This caused every audited repo to be flagged as a bulk upload. The fix was removing the flag.

## 7.3 Forward Improvements

1. Scope security rule explicitly to safe\_tool\_engineering criterion
2. Add end-to-end integration test with a real small repo
3. Extend theoretical depth detection term list to close the 4/5 gap

Final Report - B1W2 The Automaton Auditor - 10 Academy AI Mastery Program

GitHub: <https://github.com/Zerubabel-J/autonomous-code-governance-swarm-langgraph>

LangSmith Project: automaton-auditor